



RUB

RUHR-UNIVERSITÄT BOCHUM

Making the Training Visible

Computer Vision Deep Learning Project

Done by: Aya Altamimi, Simon Friz, Hyovin Kwak, Dimitri Tarnavski, Arslan Gabdulkhakov

Introduction

□ Why DeConvnet

□ What happens inside the network?

- Sparse knowledge about deeper layers

□ Go backwards through the network

- Visualization technique based on Deconvnet

□ Deconvnet


Network architecture and training

Model

MNIST: 60000 images, 6000 in each class

RPS Datasets 2520, 840 in each class

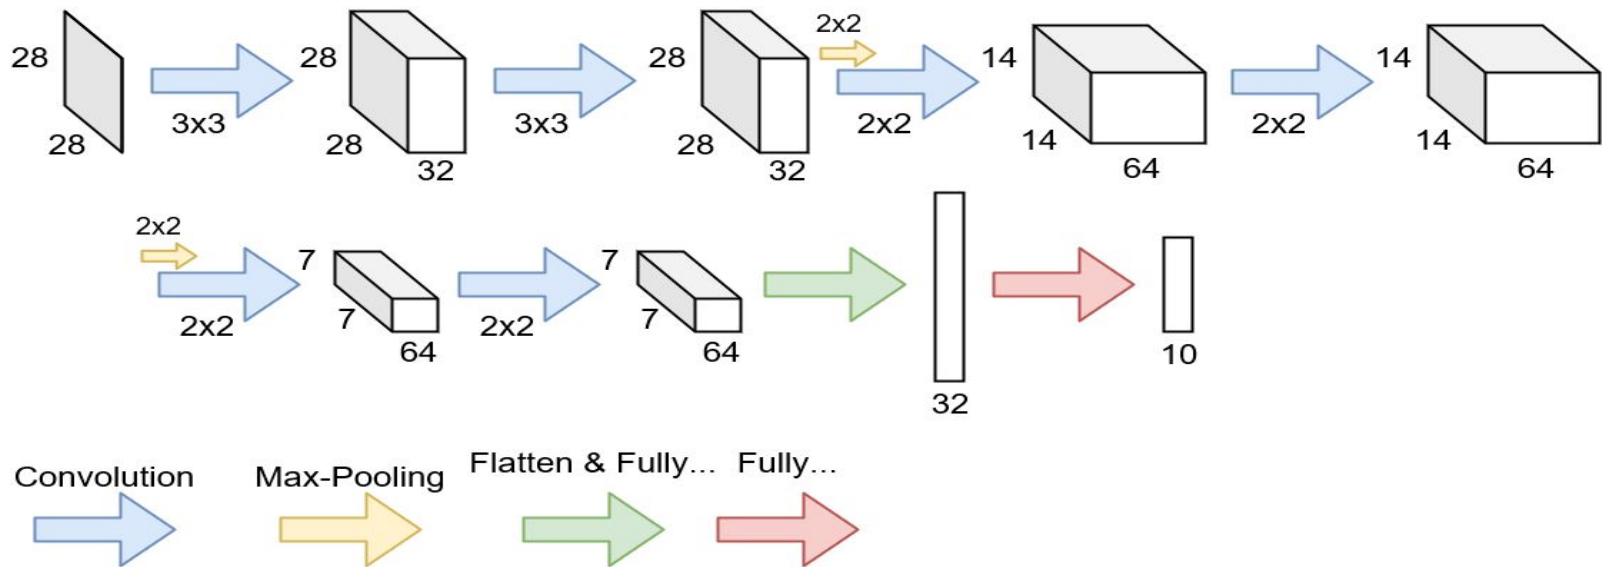
- Input (28x28 or 300x300)
- Conv2D layer (32 filters 3x3)
- Conv2D layer (64 filters 3x3)
- Max pooling (2x2)
- Dense (64)
- Output (10 or 3)



Network architecture and training

Deeper model

- Input (28x28 or 300x300)
- Conv2D layer (32 filters 3x3)
- Conv2D layer (64 filters 3x3)
- Max pooling (2x2)
- Dense (32)
- Output (10 or 3)



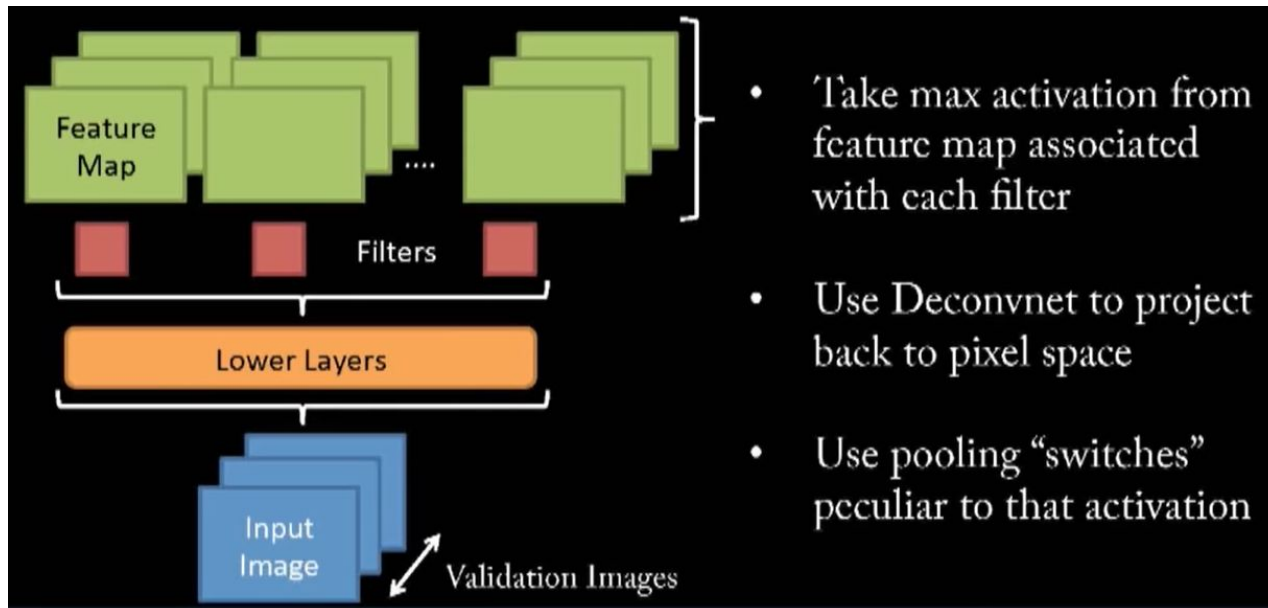
DeconvNet

➤ Idea:

- **Mapping** activities **back** to the input **pixel space**.
(Instead of mapping pixels to features, do it **the other way around**)
- **Reconstruct** the activity in the **layer beneath** that gave rise to the chosen activation.
(Reconstruct an approximate version of the convnet features from the layer beneath)
- **Repeated** until input pixel space is reached.
(DeconvNet **attached to each layer**)
- **Not** used in any **learning** capacity. (Deconvnet is only used to project back into pixel space).

Projecting back from higher layers

- **Run** input image through convolutional network (**Feed forward**)
- Results in **activations** over feature maps in **all the layers**
- Select a single feature map.
- Take strongest activation in that feature map of input image
 - Gives an idea what in the input image that activates the selected FM strongly
- Use that feature map as input of the DeconvNet



DeconvNet: unpooling \Rightarrow rectification \Rightarrow filtering

To keep edge information (avoid blurriness and color variation)

Transposed convolution

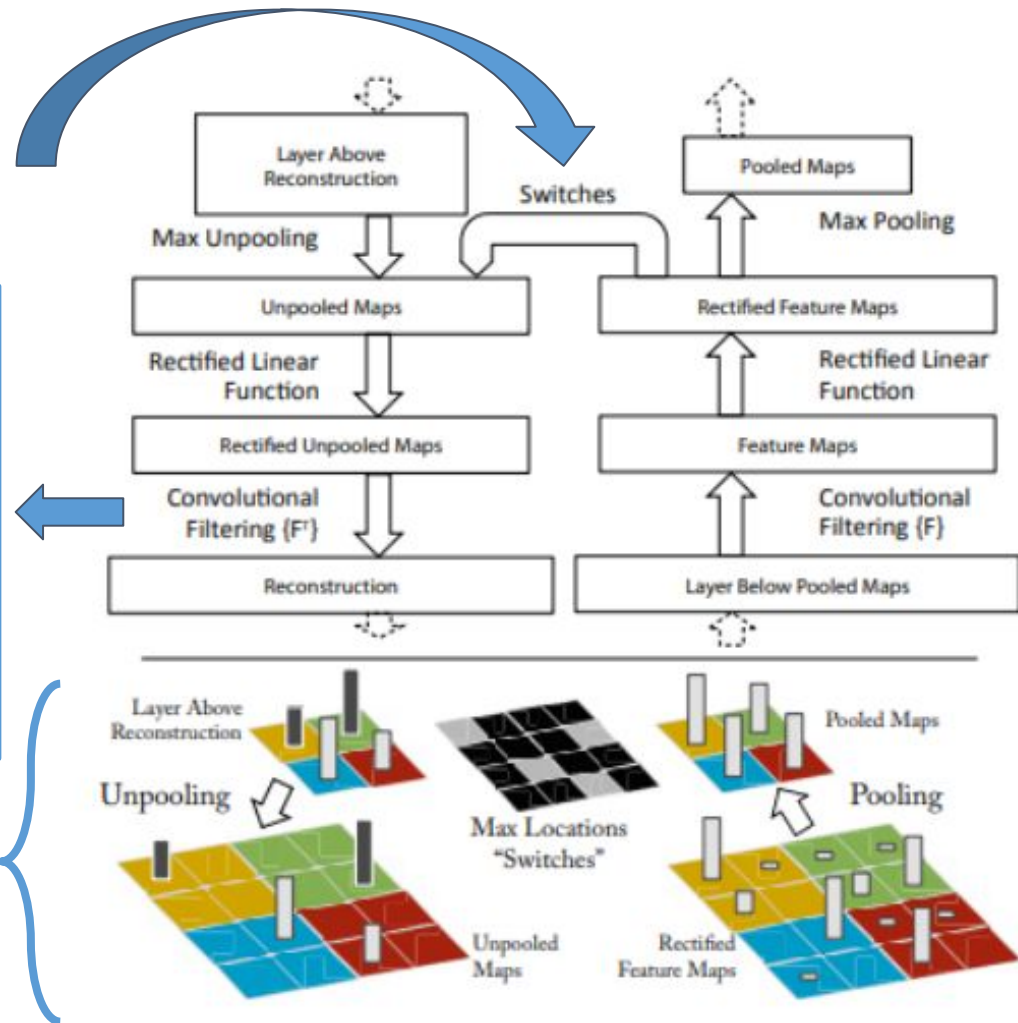
1	1	1
1	2	1
1	1	1

*


4	3
2	0

4	4	4	3	3	3
4	Ⓢ	4	3	6	3
4	4	4	3	3	3
2	2	2	0	0	0
2	4	2	0	0	0
2	2	2	0	0	0


Reversible Max pooling,
to get reconstruction clean



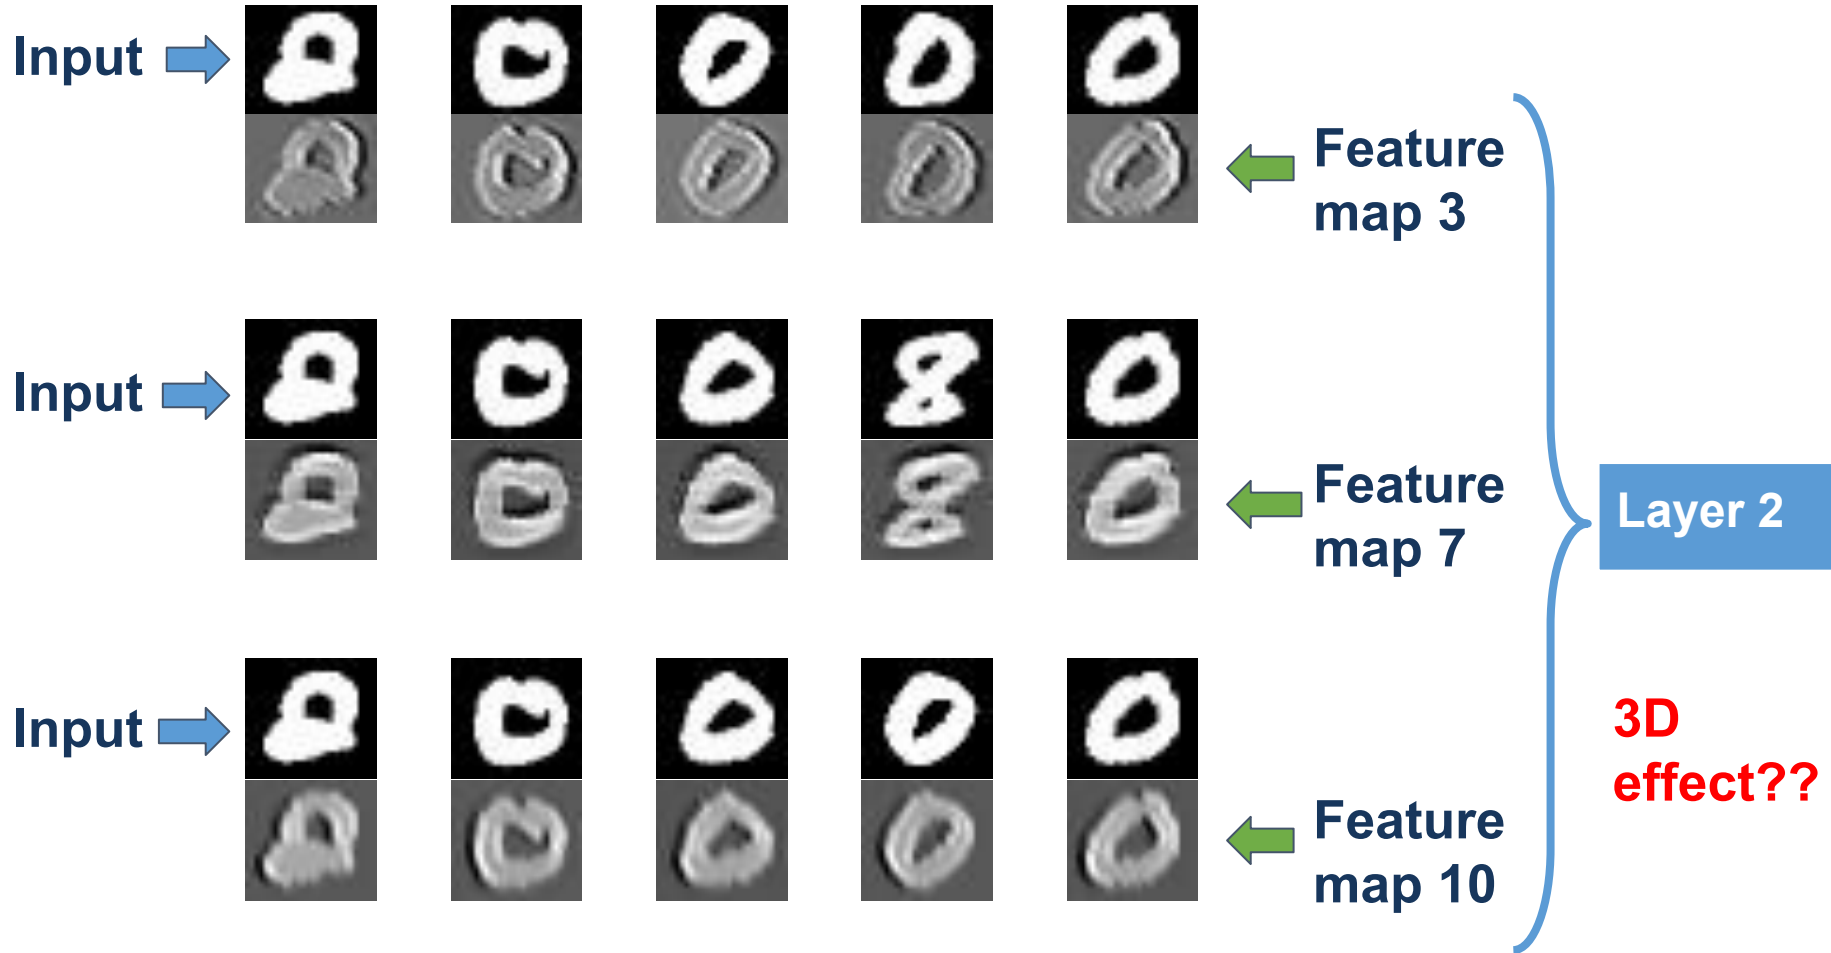
Outcomes and results



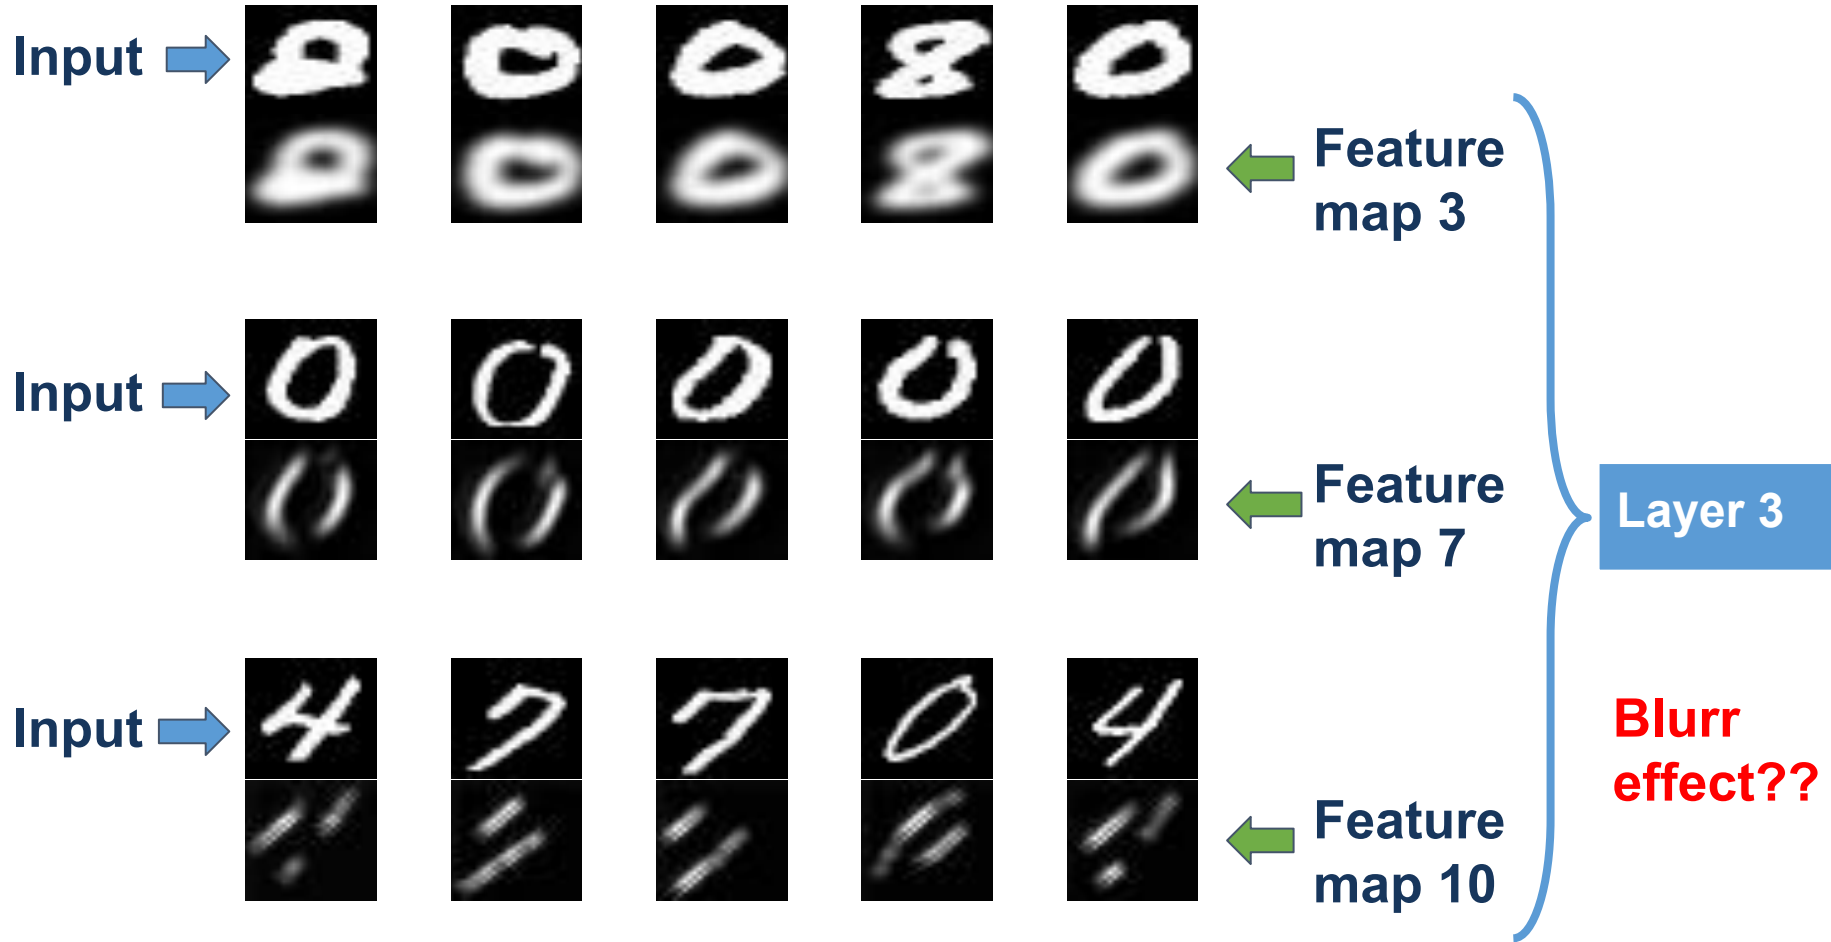
Visualisation during training (Layer 2, Feature Map 3):



Visualization of top 5 activations




Visualization of top 5 activations



Top 5 activations on a deeper model

3D effect??



Top 5 activations on a deeper model

Clumping effect??



Layer 4
FM 3




Layer 4
FM 7



Layer 5
FM 3




Layer 5
FM 7



Layer 5
FM 10


Crystal effect??

Visualization of top 5 activations for RPS dataset



Sensitivity Analysis

- **Sensitivity analysis** of the classifier output by **occluding part of the input image**.
- **Drop in classification confidence** after occlusion.



```
60000/60000 [=====] - 170s 3ms/sample - loss: 0.0051 - accuracy: 0.9994 - val_loss: 0.1577 - val_accuracy: 0.9924
current directory: /home/hkwak/Documents/Workspace/cnn-make-training-visible/curtiz_MNIST_deconv_model3_2020-01-26_16:44:33
Saved trained model at /home/hkwak/Documents/Workspace/cnn-make-training-visible/curtiz_MNIST_deconv_model3_2020-01-26_16:44:33/saved_models
model_func1 prediction: normal, blacked
model_func1.predict(normal_img): [[0. 0. 0. 0. 1. 0. 0. 0. 0. 0.]]
model_func1.predict(img_blacked): [[1.7228854e-15 1.5951092e-14 6.5399600e-12 1.5834305e-28 3.8626182e-03
0.0000000e+00 0.0000000e+00 9.9613732e-01 1.6871501e-26 4.7516211e-28]]
10000/10000 [=====] - 3s 250us/sample - loss: 0.1577 - accuracy: 0.9924
```

Reference:

Matthew D Zeiler, Rob Fergus, Visualizing and Understanding Convolutional Networks (2013)

Our Project on GitHub:



Questions?